


PREMIUM
Google search



AVISEK LAHIRI, PhD
Research @ Google | Ex-Facebook,
Twitter, Google

[+ Follow](#)

[View full profile](#)

[Please share with your networks]...

Any GOOD company will NOT be interested to start probing your knowledge on deep learning [nothing much theoretical to ask]

Be prepared with traditional concepts of

1. Bayesian min error/risk framework
2. Bias-Variance tradeoff [work out the math, not just story]
3. Naive Bayes
4. Gaussian Mixture models
5. Ensemble methods
6. Decision trees
7. Logistic regression
8. k-nearest neighbor
9. Dimensionality reduction [PCA, LDA]

Majority of times if you can't impress with traditional approaches, it is ALMOST sure you will NOT be clearing the expectation.

Even for Deep learning, have fundamental understanding of

1. How fully connected layers are different from conv layers
2. The math behind momentum, RMSProp optimizers
3. Math behind gradient explosion/vanishing
4. Why Residual Nets work as they do
5. Why batch normalization works

In summary keep your fundamentals strong, and just rely on empirical discussions/implementations.

Interested in discussions like these?

Follow [AVISEK LAHIRI, PhD](#) as we together navigate the landscapes of opportunities for internship, jobs, career, life !!